

# BIOGRAPHX

## 1. Executive Summary

### Project Title

**Graph-Augmented Agentic AI for Explainable Biomedical Question Answering**

### One-liner

Build a biomedical question answering (QA) system that doesn't just answer, but **shows its work** using:

- A **Knowledge Graph (Neo4j)**
- A **Vector Store (ChromaDB)** over PubMed-like abstracts
- A **fine-tuned open-source HuggingFace QA model**
- A **multi-agent pipeline (LangGraph)** coordinating 6 agents

**Data** comes entirely from **Kaggle**.

**Models** are **open-source** and **fine-tuned by you** (no external LLM APIs).

---

## 2. Objectives & Success Criteria

### 2.1 Main Objective

Create an **Explainable Biomedical QA system** that:

- Answers natural language medical questions
- Uses **graph reasoning** plus **retrieval-augmented generation (RAG)**

- Provides **evidence sentences and graph paths** supporting the answer

## 2.2 Success Criteria

The project is successful if:

- System answers questions end-to-end via a **Streamlit UI**
  - Each answer includes **at least one evidence sentence with PMID**
  - A **Neo4j subgraph** (e.g., Drug → TREATS → Disease) is shown for relevant questions
  - QA model achieves reasonable metrics (e.g., EM/F1) on a subset of **PubMedQA**
  - Code is reproducible and documented well enough for another student to run
- 

## 3. Problem & Business Understanding

### 3.1 Motivation

Biomedical professionals and students face:

- Huge volume of literature (PubMed, clinical studies)
- Fragmented knowledge about diseases, drugs, symptoms
- Black-box AI models that don't explain why they answered something

In medicine, **wrong but confident answers are dangerous**. We need systems that **both answer and explain**.

### 3.2 Business / Real-World Use

Potential use cases:

- Clinical decision support (junior doctors check treatments)
- Research assistants (find relationships between drugs and diseases)
- Educational tools (explain reasoning paths to students)

The project is a **prototype** for such systems, focusing on transparency, graphs, and agents.

---

## 4. Data Understanding (Kaggle Only)

You will use **three main Kaggle datasets**:

### 4.1 MedQuAD – Medical Question Answer Dataset

- Q&A pairs from NIH websites
- Used for:
  - Training/fine-tuning QA model
  - Building graph nodes (disease, drug, symptom, etc.)

### 4.2 PubMed Abstracts Dataset

- Large collection of PubMed-like abstracts
- Used as:
  - Evidence corpus for retrieval
  - Source of sentences to support final answers

#### 4.3 PubMedQA

- Biomedical QA benchmark
- Used for **evaluation only**
- Questions + associated abstracts + gold answers

#### 4.4 ETL Overview

ETL (Extract → Transform → Load) converts raw Kaggle datasets into **clean, structured, machine-ready files** stored in `data/processed/`.

##### ETL Inputs (raw data from Kaggle)

Stored in:

`data/raw/`

```
|  
|   medquad.csv  
|   pubmed_abstracts.csv  
|   pubmed_qa_pga_labeled.parquet  
|   pubmed_qa_pga_artificial.parquet
```

## ETL Outputs

Stored in:

data/processed/

```
|  
|   └── medquad_clean.csv  
|  
|   └── pubmedqa_clean.csv  
|  
|   └── pubmed_sentences.parquet  
|  
|   └── entity_mappings.csv  
└── graph/  
    |  
    |   └── nodes_disease.csv  
    |  
    |   └── nodes_drug.csv  
    |  
    |   └── nodes_question.csv  
    └── rels_about.csv
```

## ETL Pipeline Steps

### 1) clean\_medquad.py

Input: [medquad.csv](#)

Output: [medquad\\_clean.csv](#)

Purpose: Clean Q/A text → structure → remove noise

### 2) clean\_pubmedqa.py

Input: PubMedQA parquet files

Output: [pubmedqa\\_clean.csv](#)

Purpose: Normalize PubMedQA evaluation dataset

### 3) prepare\_pubmed\_sentences.py

Input: PubMed abstracts

Output: `pubmed_sentences.parquet`

Purpose: Split abstracts → sentences → store efficiently for ChromaDB

#### 4) `extract_entities.py`

Input: `medquad_clean.csv`

Output: `entity_mappings.csv`

Purpose: SciSpaCy biomedical NER → extract Disease/Drug/Symptom

#### 5) `build_graph_csvs.py`

Input: entity\_mappings + MedQuAD

Output: Neo4j node & relationship CSVs

Purpose: Build graph ingestion files

---

## 5. Tools & Technologies

### 5.1 Data & Core Libraries

- `pandas`, `numpy`, `scipy`, `tqdm`
- `matplotlib`, `seaborn` (for simple plots & EDA)

### 5.2 NLP & Biomedical NER

- `transformers`, `sentence-transformers` (HuggingFace)
- `datasets` (HuggingFace Datasets)
- `spacy`, `scispacy` (for biomedical NER)

### 5.3 Graph Tools

- `Neo4j` (Docker or Desktop)
- `py2neo` (Python driver)
- `networkx` (for graph manipulation/visualization)

## 5.4 Vector Store & Retrieval

- `chromadb`
- `faiss-cpu` (embedded within/under Chroma)

## 5.5 Agentic Framework / RAG Orchestration

- `langchain`
- `langgraph`

## 5.6 Model Fine-Tuning

- `torch`, `accelerate`
- `peft` (for LoRA)
- `bitsandbytes` (optional for 8-bit/4-bit finetuning)

## 5.7 UI & Visualization

- `streamlit`
- `pyvis` or `networkx` for graph visualization

## 5.8 Testing & Evaluation

- `pytest`, `pytest-cov`
  - `evaluate`, `rouge-score`
- 

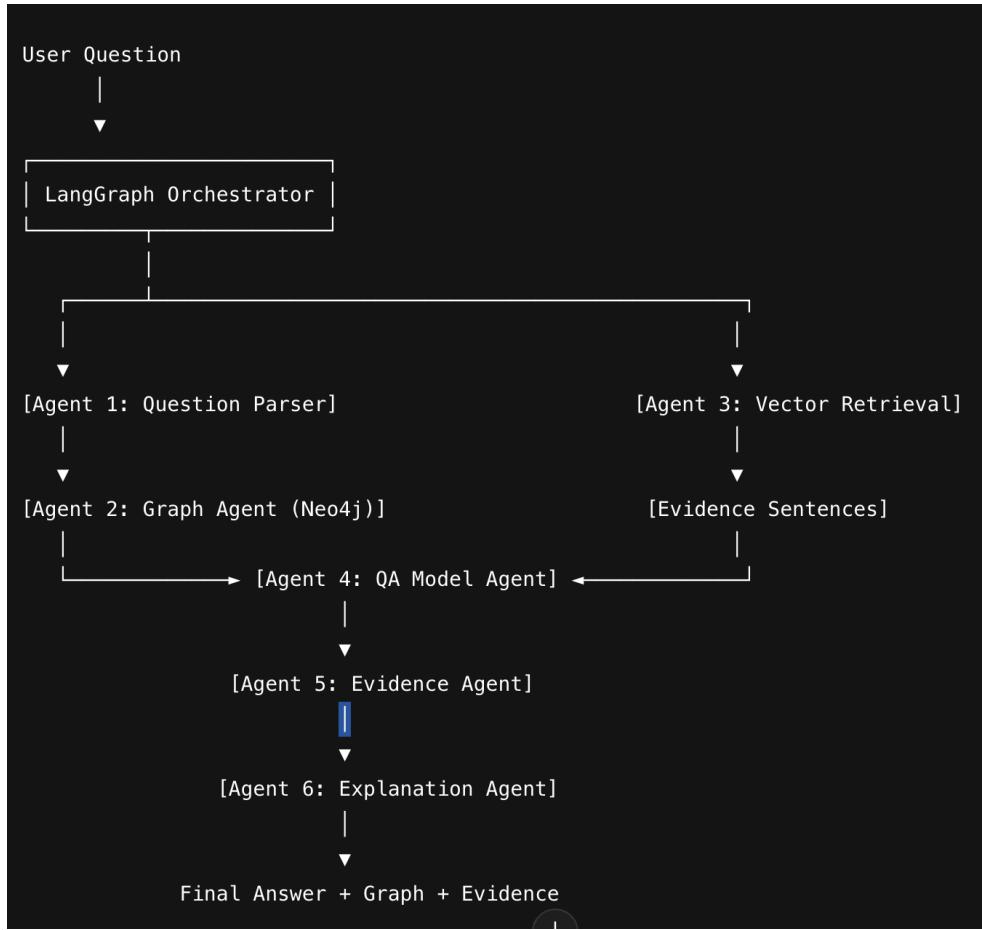
## 6. System Architecture (High-Level)

### 6.1 Big Picture

1. User asks a **biomedical question** in Streamlit.
2. **LangGraph** orchestrates 6 agents:
  - Question Parsing Agent
  - Graph Query Agent (Neo4j)
  - Vector Retrieval Agent (ChromaDB)
  - QA Model Agent (fine-tuned HF model)
  - Evidence Validation Agent
  - Explanation Agent
3. The system returns:
  - Answer
  - Evidence sentences (with PMIDs)

- Graph path(s) showing relationships

## 6.2 Architecture Diagram (Logical)




---

## 7. Agentic Workflow (6 Agents – Beginners Can Understand)

### Agent 1 – Question Parsing & Classification

- Input: user's question string

- Responsibilities:
    - Detect question type (e.g., “treatment”, “symptoms”, “cause”)
    - Extract biomedical entities (disease names, drugs, etc.) using SciSpaCy
  - Output:
    - `{"question_type": "...", "entities": [...]}`
- 

## Agent 2 – Graph Query Agent (Neo4j + Cypher)

- Input: entities + question\_type
  - Responsibilities:
    - Generate Cypher queries to explore knowledge graph
    - Fetch related nodes and edges for those entities
    - Summarize graph info (e.g., “Metformin TREATS Type 2 Diabetes”)
  - Output:
    - List of `triples` or `edges` like `(node1, relation, node2)`
- 

## Agent 3 – Vector Retrieval Agent (ChromaDB)

- Input: question + (optional) entities
  - Responsibilities:
    - Encode question → embedding (sentence-transformers)
    - Retrieve top-k similar PubMed sentences
    - Attach PMIDs and metadata
  - Output:
    - `[{"pmid": ..., "sentence": "...", "score": ...}, ...]`
- 

#### Agent 4 – QA Model Agent (Fine-tuned HF Model)

- Input: question + graph summary + top evidence sentences
  - Responsibilities:
    - Construct a prompt/input (e.g. for FLAN-T5)
    - Generate an answer sequence
  - Output:
    - Answer text (e.g., “Metformin and insulin are common treatments for type 2 diabetes.”)
- 

#### Agent 5 – Evidence Validation Agent

- Input: answer + evidence sentences
  - Responsibilities:
    - Check which sentences strongly support the answer
    - Optionally, use a simple similarity scoring between answer and each sentence
  - Output:
    - Ranked evidence; top-N sentences retained
- 

## Agent 6 – Explanation & Output Agent

- Input: answer + graph triples + evidence list
- Responsibilities:
  - Build structured final output:
    - Explanation text if needed
    - Evidence list with PMIDs
    - Graph paths (e.g., drug–disease–symptom)
- Output (JSON-like):

{

"answer": "Metformin is commonly used to treat type 2 diabetes.",

```
"graph_path": [  
    "Metformin", "TREATS", "Type 2 Diabetes"]  
,  
"evidence": [  
    {"pmid": "12345", "sentence": "Metformin is a first-line therapy ..."}  
]  
}
```

---

## 8. Repository Structure (Project Layout)

```
BIOGRAPHX/  
|  
|   └── README.md  
|  
|   └── requirements.txt  
|  
|   └── .gitignore  
|  
|  
|   └── data/  
|       |   └── raw/  
|       |       └── medquad.csv  
|       |       └── pubmed_abstracts.csv  
|       |       └── pubmed_qa_pga_labeled.parquet
```

```
|   |   └── pubmed_qa_pga_artificial.parquet  
|  
|  
|   └── processed/  
|       ├── medquad_clean.csv  
|       ├── pubmedqa_clean.csv  
|       ├── pubmed_sentences.parquet  
|       ├── entity_mappings.csv  
|       └── graph/  
|           ├── nodes_disease.csv  
|           ├── nodes_drug.csv  
|           ├── nodes_symptom.csv (optional)  
|           ├── nodes_question.csv  
|           └── rels_about.csv  
  
|  
└── etl/  
    ├── clean_medquad.py  
    ├── clean_pubmedqa.py  
    ├── prepare_pubmed_sentences.py  
    ├── extract_entities.py  
    ├── build_graph_csvs.py  
    └── utils/
```

```
|   └── __init__.py
|   └── text_cleaning.py
|
|   └── graph/
|       ├── schema.cql
|       ├── load_graph.py
|       ├── extract_entities.py (if used separately)
|       └── queries/
|           ├── get_disease_relations.cql
|           ├── get_drug_relations.cql
|           └── get_question_entities.cql
|
|   └── rag/
|       ├── prepare_pubmed_sentences.py (duplicate optional)
|       ├── build_index.py
|       ├── retriever.py
|       └── utils/
|           └── embedding_models.py
|
|   └── agents/
|       └── question_agent.py
```

```
|   ├── graph_agent.py  
|   ├── retriever_agent.py  
|   ├── qa_model_agent.py  
|   ├── evidence_agent.py  
|   ├── explanation_agent.py  
|   └── agent_graph.py  
  
|  
└── training/  
    ├── prepare_medquad_dataset.py  
    ├── finetune_qa_model.py  
    ├── test_finetuned_model.py  
    ├── eval_pubmedqa.py  
    └── utils/  
        └── tokenizer_utils.py  
  
|  
└── models/  
    └── fine_tuned/  
        └── flan_t5_medquad/  
            ├── config.json  
            ├── adapter_model.bin  
            └── tokenizer.json
```

```
|      └── special_tokens_map.json  
|  
|  
└── app/  
    |   ├── streamlit_app.py  
    |   └── components/  
    |       ├── graph_viz.py  
    |       └── evidence_panel.py  
|  
|  
└── configs/  
    |   ├── config.yaml  
    |   ├── neo4j.yaml  
    |   ├── chroma.yaml  
    |   └── training.yaml  
|  
|  
└── notebooks/  
    |   ├── 01_eda_medquad.ipynb  
    |   ├── 02_pubmed_analysis.ipynb  
    |   ├── 03_error_analysis.ipynb  
    |   └── 04_agent_pipeline_debug.ipynb  
|  
|  
└── tests/
```

```
|   └── test_data_loading.py  
|   └── test_graph_build.py  
|   └── test_retriever.py  
|   └── test_agents.py  
|       └── test_pipeline_integration.py  
|  
└── scripts/  
    ├── run_neo4j_docker.sh  
    ├── build_all_etl.sh  
    ├── run_all_tests.sh  
    └── demo_pipeline.py
```

---

## 9. Environment & Setup (Beginner-Friendly)

### 9.1 Create Project & Virtual Environment

```
mkdir graph-agentic-biomed-qa
```

```
cd graph-agentic-biomed-qa
```

```
python -m venv .venv
```

```
# Activate
```

```
# Windows:  
.\venv\Scripts\activate  
  
# macOS/Linux:  
source .venv/bin/activate
```

## 9.2 Create requirements.txt and Install

Use the dependency list you already have (with datasets, peft, bitsandbytes, etc.) and run:

```
pip install -r requirements.txt
```

## 9.3 Initialize Git

```
git init  
  
echo ".venv/" >> .gitignore  
  
echo "__pycache__/" >> .gitignore  
  
echo "data/raw/" >> .gitignore
```

---

# Project Implementation Plan

## Sprint 1 – Project Bootstrapping & Data Download (DETAILED)

### **Goal:**

Get the project repository, environment, dependencies, and raw datasets ready with basic exploration.

---

### **Tasks:**

#### **1. Create Repository & Virtual Environment**

- `mkdir graph-agentic-biomed-qa`
- Initialize Git repo
- Create virtual environment (`python -m venv .venv`)
- Activate venv
- Add `.venv/`, `__pycache__/`, and `data/raw/` to `.gitignore`

#### **2. Install Requirements**

- Create `requirements.txt` containing:
  - pandas, numpy, tqdm
  - scispacy, spacy
  - transformers, sentence-transformers
  - chromadb, faiss-cpu
  - py2neo, neo4j

- torch, accelerate, peft
  - streamlit, pyvis
- Install:  
`pip install -r requirements.txt`

### 3. Download Kaggle Datasets

Use Kaggle CLI or manual download:

- **MedQuAD** → `data/raw/medquad/`
- **PubMed Abstracts** → `data/raw/pubmed/`
- **PubMedQA** → `data/raw/pubmedqa/`

### 4. Create EDA Notebook

File: `notebooks/01_eda_medquad.ipynb`

Perform initial understanding:

- Load MedQuAD JSON / CSV
- Show `.head()`
- Print a couple Q/A entries
- Count rows
- Identify missing values
- Basic text length statistics

---

**Deliverables:**

- Working virtual environment
  - Raw Kaggle datasets inside `data/raw/`
  - EDA notebook demonstrating dataset familiarity
- 

## Sprint 2 – Data Cleaning, NER & Knowledge Graph Schema (FULL DETAILED VERSION)

*(All content below matches your formatting but expanded)*

### **Goal:**

Clean MedQuAD, extract biomedical entities using SciSpaCy, and design the Neo4j graph schema.

---

### **Tasks:**

---

#### **1. Clean MedQuAD – `training/prepare_medquad_dataset.py`**

##### **Actions:**

- Load raw `medquad.csv`
- Normalize text:
  - trim whitespace
  - remove HTML tags
  - unify punctuation
  - lowercase (optional)

- Keep only relevant columns:
  - `question`
  - `answer`
  - `source_url` (*optional*)
- Drop duplicates and rows where Q/A is missing
- Save cleaned dataset to:  
`data/processed/medquad_clean.csv`

#### Purpose:

This dataset becomes:

- Input to NER
- Graph Question nodes
- Training set for QA model fine-tuning

---

## 2. Install SciSpaCy + Perform Biomedical NER – `graph/extract_entities.py`

#### Actions:

Install SciSpaCy model:

```
pip install scispacy
```

```
pip install https://.../en_ner_bc5cdr_md-0.5.0.tar.gz
```

- 
- For each question + answer:
  - run SciSpaCy

- extract named entities
- Map entities to:
  - **Disease**
  - **Drug**
  - **Symptom**
  - **Chemical** (*optional*)
- Store:
  - entity\_text
  - entity\_type
  - originating\_question
- Save output as:  
`data/processed/entity_mappings.csv`

#### Purpose:

Provides **all entity nodes** for the knowledge graph.

---

### 3. Design Graph Schema – `graph/schema.cql`

#### Actions:

Define Neo4j labels:

- `:Question`
- `:Disease`

- :Drug
- :Symptom

Define relationships:

- (:Question)-[:ABOUT]->(:Disease | :Drug | :Symptom)
- (Optional) :TREATS, :CAUSES, :HAS\_SYMPTOM

Define constraints:

```
CREATE CONSTRAINT question_unique IF NOT EXISTS
```

```
ON (q:Question) ASSERT q.id IS UNIQUE;
```

```
CREATE CONSTRAINT disease_unique IF NOT EXISTS
```

```
ON (d:Disease) ASSERT d.name IS UNIQUE;
```

### **Purpose:**

Creates the structure for the BioGraphX Neo4j Knowledge Graph.

---

## **4. Build Graph CSVs – [etl/build\\_graph\\_csvs.py](#)**

### **Actions:**

- Convert entity mappings into node CSVs:
  - nodes\_disease.csv
  - nodes\_drug.csv

- `nodes_symptom.csv`
- `nodes_question.csv`
- Create ABOUT relationships CSV:
  - `rels_about.csv` linking Q → entities
- Save all to:  
`data/processed/graph/`

---

#### Purpose:

These files are consumed by Neo4j loader in Sprint 3.

---

#### Deliverables:

- `medquad_clean.csv`
  - `entity_mappings.csv`
  - `graph/schema.cql`
  - Full graph CSV files (nodes + edges)
  - Graph schema documentation in README
- 

## Sprint 3 – Neo4j Graph Build + PubMed Sentence Embeddings + Chroma Setup (FULL DETAILED VERSION)

#### Goal:

Load graph into Neo4j, prepare PubMed sentence corpus, embed sentences using BioBERT, and build Chroma vector store.

---

**Tasks:**

---

## 1. Start Neo4j DB

Choose one:

### Option A — Docker

```
docker run \
-p7474:7474 -p7687:7687 \
-d --name neo4j \
-e NEO4J_AUTH=neo4j/password \
neo4j:5.2
```

### Option B — Neo4j Desktop

- Create project
- Enable Bolt
- Enable APOC if needed

---

## 2. Load Graph – `graph/load_graph.py`

### Actions:

- Connect using py2neo
- Apply schema constraints from `schema.cql`
- Load node CSVs:

- Disease, Drug, Symptom, Question
- Load relationship CSV:
  - ABOUT edges

**Purpose:**

Builds the **Biomedical Knowledge Graph** used by Agent 2.

---

### 3. Prepare PubMed Sentence Corpus – [rag/prepare\\_pubmed\\_sentences.py](#)

**Actions:**

- Read [pubmed\\_abstracts.csv](#)
- For each abstract:
  - extract [pmid](#)
  - split abstract into individual sentences
- Store rows:

[pmid, sentence](#)

- Save to:  
[data/processed/pubmed\\_sentences.parquet](#)

**Purpose:**

Provides a retrieval-friendly dataset for embedding.

---

#### **4. Build Chroma Vector Index – [rag/build\\_index.py](#)**

##### **Actions:**

- Load PubMed sentences
- Embed using BioBERT / SciBERT:

```
model = SentenceTransformer("pritamdeka/S-BioBERT-snli-mnli-stsb")
```

- Insert embeddings + metadata into Chroma collection:  
[pubmed\\_sentences](#)
- Persist vector store to disk

##### **Purpose:**

Used by Vector Retrieval Agent (Agent 3) to fetch evidence.

---

#### **5. Implement Retrieval – [rag/retriever.py](#)**

##### **Function:**

```
retrieve_sentences(query, k=5)
```

##### **Steps:**

- Embed query sentence
- Query Chroma for top-k nearest sentences
- Return:

[

```
{"pmid": "...", "sentence": "...", "score": ...},  
...  
]
```

**Purpose:**

Provides evidence sentences for QA reasoning.

---

**Deliverables:**

- Fully populated **Neo4j DB**
  - `pubmed_sentences.parquet`
  - Chroma collection `pubmed_sentences`
  - Working retriever returning biomedical evidence
- 

## Sprint 4 – Fine-Tuning HuggingFace QA Model (DETAILED)

**Goal:**

Train a domain-specific QA model on MedQuAD.

---

**Tasks:**

### 1. Prepare HF Dataset – `training/prepare_medquad_dataset.py`

- Load `medquad_clean.csv`
- Build dataset:

```
{"input": question, "output": answer}
```

- 
- Split: train/val

## 2. Choose Model

- [google/flan-t5-base](#) (recommended)
  - Alternatives: [t5-small](#), [t5-base](#)
- 

## 3. Write Training Script – [training/finetune\\_qa\\_model.py](#)

- Tokenize inputs (max\_source\_length=512)
  - Add LoRA configuration for speed
  - Train for 2–4 epochs
  - Save model to:  
[models/fine\\_tuned/flan\\_t5\\_medquad/](#)
- 

## 4. Test Fine-Tuned Model – [training/test\\_finetuned\\_model.py](#)

- Load saved model
  - Run inference on 2–3 sample questions
  - Print outputs
-

## **Deliverables:**

- Fine-tuned model
  - Training logs
  - Test script confirming basic Q/A performance
- 

## **Sprint 5 – Implement the 6-Agent LangGraph Pipeline (DETAILED)**

### **Goal:**

Enable the full reasoning workflow using 6 specialized agents.

---

### **Tasks:**

#### **1. Implement Each Agent (agents/\*.py)**

##### **question\_agent.py**

- Input: user question text
- Output:

```
{  
    "entities": [...],  
    "question_type": ...  
}
```

##### **graph\_agent.py**

- Runs Cypher queries on Neo4j
- Retrieves graph relationships for found entities

### **retriever\_agent.py**

- Queries Chroma vector DB
- Returns evidence sentences with PMIDs

### **qa\_model\_agent.py**

- Combines:
  - question
  - evidence
  - graph context
- Calls fine-tuned T5 model

### **evidence\_agent.py**

- Scores sentences vs. model answer
- Picks top explanatory evidence

### **explanation\_agent.py**

- Final JSON output:

answer

graph\_path

evidence

---

## 2. Build LangGraph Orchestration – `agent_graph.py`

- Define node functions
- Define state dictionary
- Connect agents in order:

Question → QuestionAgent → GraphAgent →  
RetrieverAgent → QA Agent → EvidenceAgent → ExplanationAgent

---

## 3. Demo Script – `scripts/run_demo_question.py`

- Hardcode sample question
- Run pipeline
- Print structured output

---

### Deliverables:

- Six working agent modules
  - LangGraph orchestrator
  - Working end-to-end demo
-

## Sprint 6 – Streamlit UI + Evaluation + Final Polish (DETAILED)

### Goal:

Create user-friendly frontend, run evaluation, finalize documentation.

---

### Tasks:

---

#### 1. Streamlit UI – [app/streamlit\\_app.py](#)

##### UI Components:

- Input textbox for biomedical question
  - “Ask” button
  - Output panels:
    - Final Answer
    - Evidence sentences
    - Graph visualization (pyvis)
- 

#### 2. Evaluation – [training/eval\\_pubmedqa.py](#)

##### Steps:

- Load small PubMedQA sample
- For each question:
  - run your pipeline

- compare output vs. gold answer
  - Compute:
    - Exact Match (EM)
    - F1 Score
- 

### 3. Error Analysis – notebooks/02\_error\_analysis.ipynb

Inspect failed cases:

- Was evidence relevant?
  - Did graph reasoning help?
  - Did model hallucinate?
  - Are entities missing?
- 

### 4. Final Polish

Update README:

- Setup instructions
  - How to run Neo4j
  - How to run Streamlit app
  - How to run retriever
  - How to evaluate
-

**Deliverables:**

- Fully working Streamlit app
- Evaluation metrics
- Visual graph explanations
- Error analysis notebook
- Final documentation